



The University of Texas at Austin  
Oden Institute for Computational  
Engineering and Sciences

**28 April 2023**

# PARALLEL SUPPORT VECTOR MACHINES

---

All relevant code is at link: <https://github.com/V-Rang/CSE-392-Parallel-Algorithms-for-Scientific-Computing-Project--Parallel-Support-Vector-Machines>

**Dilip Geethakrishnan and Venugopal Ranganathan**

# OUTLINE

---

1. Problem Statement
2. Sequential – Complexity & Algorithm
3. Parallel – Complexity & Algorithm
4. Implementation Details
5. Results

# OUTLINE

---

- 1. Problem Statement**
2. Sequential – Complexity & Algorithm
3. Parallel – Complexity & Algorithm
4. Implementation Details
5. Results

# Problem Statement

---

$$\begin{array}{ll} \text{minimize} & f_o(\vec{w}, \vec{\xi}) = \frac{1}{2} \vec{w}^T \vec{w} + C \sum_{i=1}^n \xi_i \\ \text{s. t.} & y_i(\vec{w}^T \vec{x}_i + b) \geq 1 - \xi_i, \quad i = 1, \dots, n \\ & \xi_i \geq 0, \quad i = 1, \dots, n \end{array} \longrightarrow \begin{bmatrix} Q & \vec{1}_m \\ \vec{1}_m^T & 0 \end{bmatrix} \cdot \begin{bmatrix} \vec{\alpha} \\ b \end{bmatrix} = \begin{bmatrix} \vec{y} \\ 0 \end{bmatrix},$$

The aim is to improve/study some performance of SVM's. We do this by the following:

- Parallelize Matrix multiplications on GPUs
- Improve performance of Matrix inversion by using parallelized Randomized SVD.

In this study, we compare the performances of some standard SVM libraries, as well as modify a variant of SVM, called Least-Squared SVM. In Least squared SVM, we arrive at a step where we have a system to solve, of the form  $Ax=b$ . We then use the RSVD to find the pseudo inverse.

# OUTLINE

---

1. Problem Statement
- 2. Sequential – Complexity & Algorithm**
3. Parallel – Complexity & Algorithm
4. Implementation Details
5. Results

# Sequential – Complexity & Algorithm

Matrix-matrix multiplication –  $O(n^3)$

---

## Algorithm 1 Naive matrix-matrix multiplication

---

**Require:**  $A_{m,p}, B_{p,n}$

```
1: Initialize  $C$  to be a zero matrix of size  $m \times n$ 
2: for  $i = 1, \dots, m$  do
3:   for  $j = 1, \dots, n$  do
4:     for  $k = 1, \dots, p$  do
5:        $C_{ij} = C_{ij} + A_{ik}B_{kj}$ 
6:     end for
7:   end for
8: end for
9: Return  $C$ 
```

---

Randomized Singular Value Decomposition –  $O(mnk + nk^2)$

### ALGORITHM: RSVD — BASIC RANDOMIZED SVD

*Inputs:* An  $m \times n$  matrix  $\mathbf{A}$ , a target rank  $k$ , and an over-sampling parameter  $p$  (say  $p = 10$ ).

*Outputs:* Matrices  $\mathbf{U}$ ,  $\mathbf{D}$ , and  $\mathbf{V}$  in an approximate rank- $(k + p)$  SVD of  $\mathbf{A}$  (so that  $\mathbf{U}$  and  $\mathbf{V}$  are orthonormal,  $\mathbf{D}$  is diagonal, and  $\mathbf{A} \approx \mathbf{UDV}^*$ .)

**Stage A:**

(1) Form an  $n \times (k + p)$  Gaussian random matrix  $\mathbf{G}$ .

$\mathbf{G} = \text{randn}(n, k+p)$

(2) Form the sample matrix  $\mathbf{Y} = \mathbf{A}\mathbf{G}$ .

$\mathbf{Y} = \mathbf{A} * \mathbf{G}$

(3) Orthonormalize the columns of the sample matrix  $\mathbf{Q} = \text{orth}(\mathbf{Y})$ .

$[\mathbf{Q}, \sim] = \text{qr}(\mathbf{Y}, 0)$

**Stage B:**

(4) Form the  $(k + p) \times n$  matrix  $\mathbf{B} = \mathbf{Q}^* \mathbf{A}$ .

$\mathbf{B} = \mathbf{Q}' * \mathbf{A}$

(5) Form the SVD of the small matrix  $\mathbf{B}$ :  $\mathbf{B} = \hat{\mathbf{U}}\mathbf{D}\mathbf{V}^*$ .

$[\mathbf{Uhat}, \mathbf{D}, \mathbf{V}] = \text{svd}(\mathbf{B}, 'econ')$

(6) Form  $\mathbf{U} = \mathbf{Q}\hat{\mathbf{U}}$ .

$\mathbf{U} = \mathbf{Q} * \mathbf{Uhat}$

“Finding Structure with Randomness” -  
Halko, Martinsson, Tropp (2011)

# OUTLINE

---

1. Problem Statement
2. Sequential – Complexity & Algorithm
- 3. Parallel – Complexity & Algorithm**
4. Implementation Details
5. Results

## Parallel – Complexity & Algorithm

---

Naïve Gpu Matrix-matrix multiplication shown for analysis.

Work:  $O(mn^2)$ .

Depth:  $O(n)$

---

### Algorithm 2 Naive GPU matmul

---

Require:  $A_{m,p}, B_{p,n}$

Initialize  $C$  to be a zero matrix of size  $m \times p$

2: parfor  $i = 1, \dots, m$  do

    parfor  $j = 1, \dots, p$  do

4:        for  $k = 1, \dots, n$  do

$C_{ij} = C_{ij} + A_{ik}B_{kj}$

6:        end for

    end parfor

8: end parfor

Return  $C$

---



## Parallel – Complexity & Algorithm

---

Naïve Gpu Matrix-matrix multiplication shown for analysis. We replace the other operations with parallel variants from various libraries and compare results.

For simplicity, we assume  $m=n$

Steps:

Forming omega: work =  $O(n*k)$  Depth =  $O(n*k)$

$Y=A*\Omega$  ,  $O(nk^2)$ ,  $O(k)$

$Q=QR(Y)$   $O(nk^2)$

$B=Q^T A$   $O(nk^2)$  ,  $O(k)$

$U'DVh=SVD(B)$   $O(nk^2)$

$U=QU'$   $O(nk^2)$  ,  $O(k)$

---

### Algorithm 3 Simplified parallel RSVD

---

**Require:**  $A_{m,n}$  With approx rank= $k$

Form  $\Omega_{n,k}$ , a random matrix.

$Y=Naive\_GPU\_matmul(A,\Omega)$

3:  $Q,-=parallel\_QR(Y)$

$B=Naive\_GPU\_matmul(Q^T,A)$

$U'DVh=ParallelSVD(B)$

6:  $U=Naive\_GPU\_matmul(Q,U')$

Return  $U, D, Vh$

---

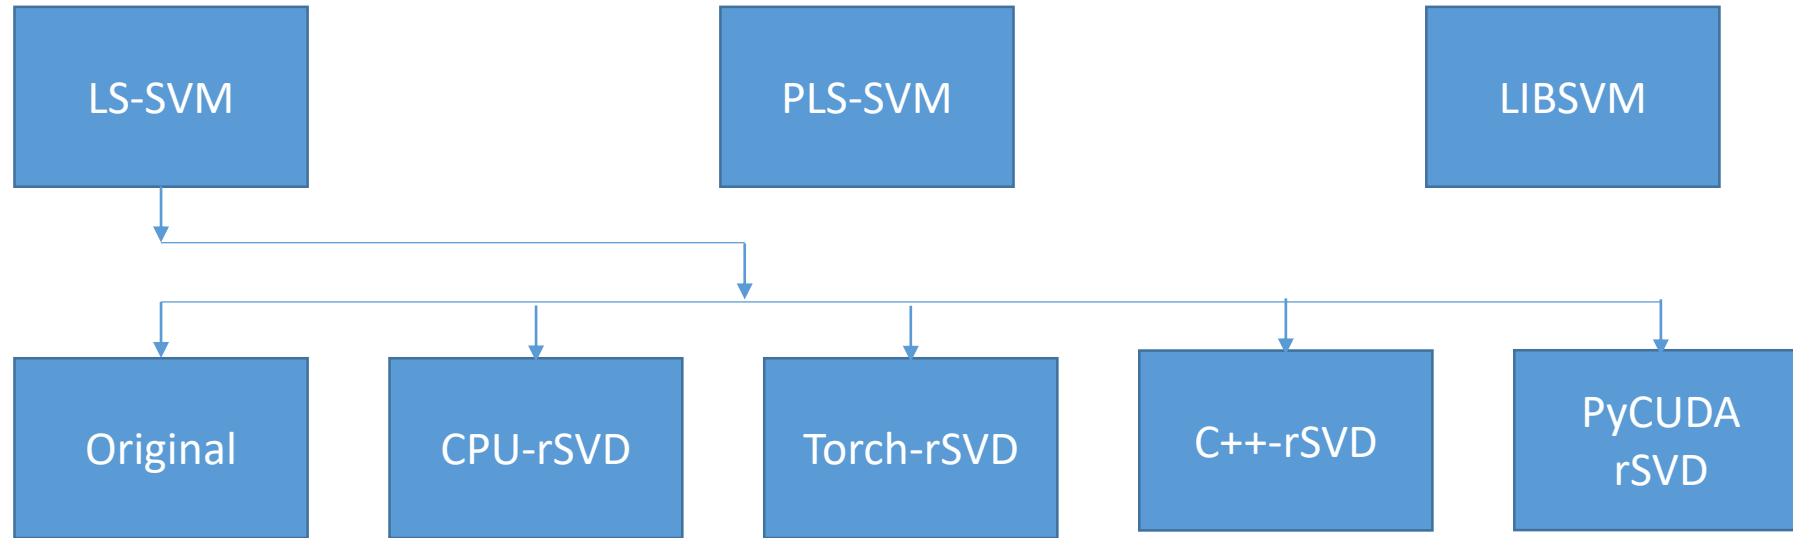
# OUTLINE

---

1. Problem Statement
2. Sequential – Complexity & Algorithm
3. Parallel – Complexity & Algorithm
- 4. Implementation Details**
5. Results

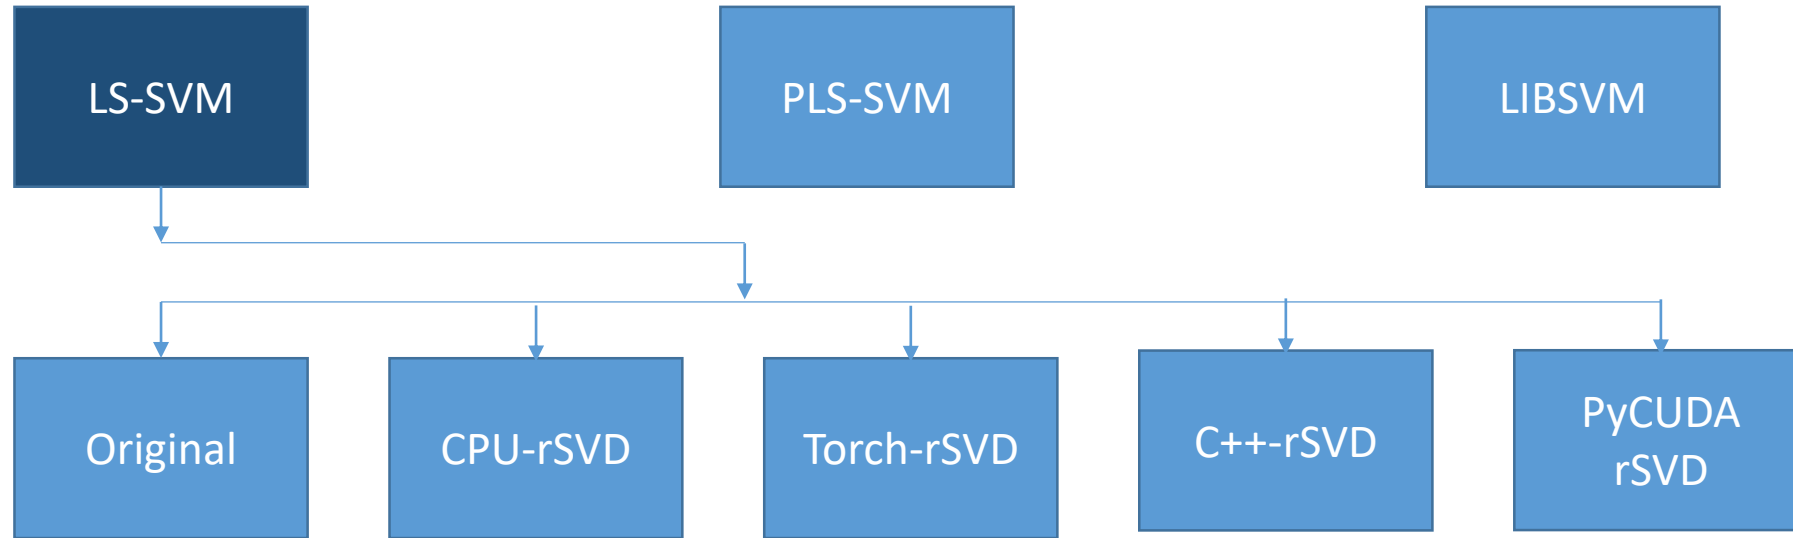
# Implementation Details

---




# Implementation Details

---



# Implementation Details




$$\begin{aligned} \text{minimize} \quad & f_o(\vec{w}, \vec{\xi}) = \frac{1}{2} \vec{w}^T \vec{w} + C \sum_{i=1}^n \xi_i \\ \text{s.t.} \quad & y_i(\vec{w}^T \vec{x}_i + b) \geq 1 - \xi_i, \quad i = 1, \dots, n \\ & \xi_i \geq 0, \quad i = 1, \dots, n \end{aligned}$$

$$\begin{aligned} \text{minimize} \quad & f_o(\vec{w}, \vec{\xi}) = \frac{1}{2} \vec{w}^T \vec{w} + \gamma \frac{1}{2} \sum_{i=1}^n \xi_i^2 \\ \text{s.t.} \quad & y_i(\vec{w}^T \vec{x}_i + b) = 1 - \xi_i, \quad i = 1, \dots, n \end{aligned}$$

$$\begin{bmatrix} Q & \vec{1}_m \\ \vec{1}_m^T & 0 \end{bmatrix} \cdot \begin{bmatrix} \vec{\alpha} \\ b \end{bmatrix} = \begin{bmatrix} \vec{y} \\ 0 \end{bmatrix},$$

# Implementation Details



```
def fit(self, X, y):
    """Fits the model given the set of X attribute vectors and y labels.
    - X: ndarray of shape (n_samples, n_attributes)
    - y: ndarray of shape (n_samples,) or (n_samples, n)
        If the label is represented by an array of n elements, the y
        parameter must have n columns.
    """
    y_resaped = y.reshape(-1,1) if y.ndim==1 else y

    self.sv_x = X
    self.sv_y = y_resaped
    self.y_labels = np.unique(y_resaped, axis=0)


    if len(self.y_labels) == 2: # binary classification
        # converting to -1/+1
        y_values = np.where(
            (y_resaped == self.y_labels[0]).all(axis=1)
            , -1, +1)[:, np.newaxis] # making it a column vector

        self.b, self.alpha = self._optimize_parameters(X, y_values)
```

```
def _optimize_parameters(self, X, y_values):
    """Help function that optimizes the dual variables through the
    use of the kernel matrix pseudo-inverse.
    """
    sigma = np.multiply(y_values*y_values.T, self.K(X,X))
    #print(y_values.shape)
    # print(self.K(X,X).shape)
    A = np.block([
        [0, y_values.T],
        [y_values, sigma + self.gamma**-1 * np.eye(len(y_values))]
    ])
    B = np.array([0]+[1]*len(y_values))
```

# Implementation Details

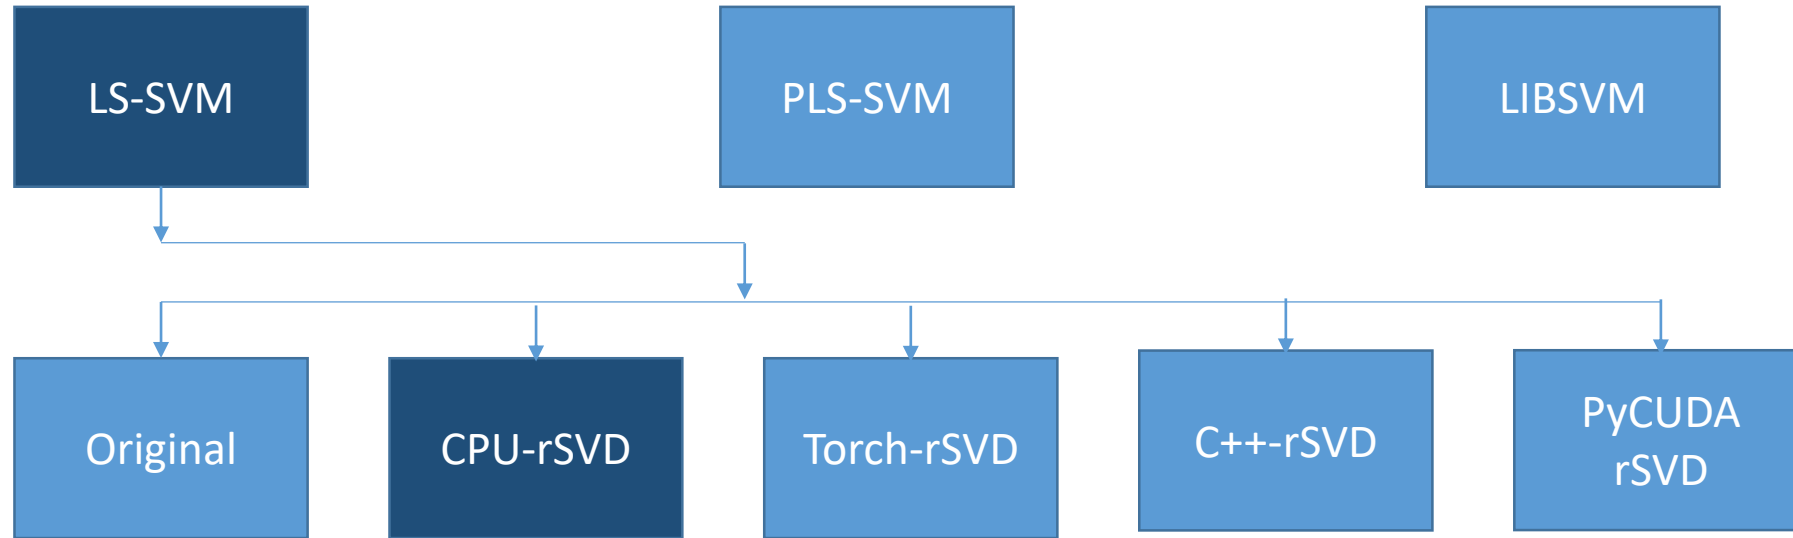
---



```
if (use_cpu_original):  
    print('I Doing CPU')  
    U,S,V = np.linalg.svd(A)  
    A_cross = V.T@np.linalg.inv(np.diag(S))@U.T  
    solution = dot(A_cross, B)
```

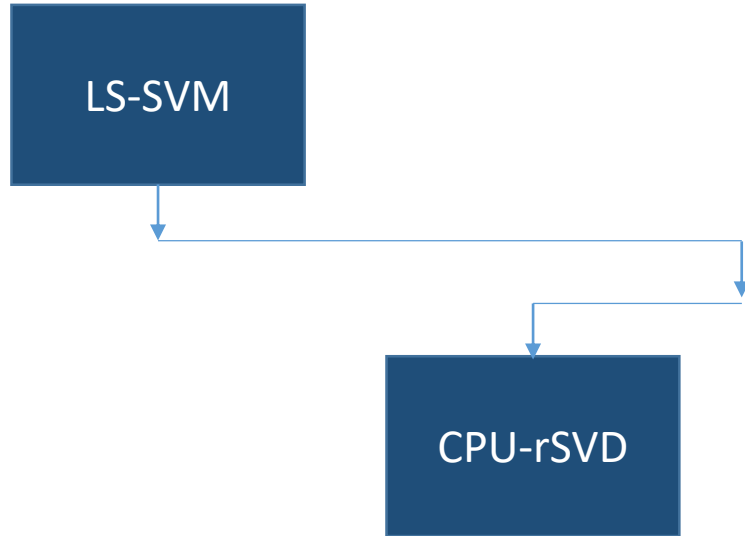
# Implementation Details

---





# Implementation Details



## ALGORITHM: RSVD — BASIC RANDOMIZED SVD

*Inputs:* An  $m \times n$  matrix  $\mathbf{A}$ , a target rank  $k$ , and an over-sampling parameter  $p$  (say  $p = 10$ ).

*Outputs:* Matrices  $\mathbf{U}$ ,  $\mathbf{D}$ , and  $\mathbf{V}$  in an approximate rank- $(k + p)$  SVD of  $\mathbf{A}$  (so that  $\mathbf{U}$  and  $\mathbf{V}$  are orthonormal,  $\mathbf{D}$  is diagonal, and  $\mathbf{A} \approx \mathbf{U}\mathbf{D}\mathbf{V}^*$ .)

### Stage A:

(1) Form an  $n \times (k + p)$  Gaussian random matrix  $\mathbf{G}$ .

$\mathbf{G} = \text{randn}(n, k+p)$

(2) Form the sample matrix  $\mathbf{Y} = \mathbf{A}\mathbf{G}$ .

$\mathbf{Y} = \mathbf{A} * \mathbf{G}$

(3) Orthonormalize the columns of the sample matrix  $\mathbf{Q} = \text{orth}(\mathbf{Y})$ .

$[\mathbf{Q}, \sim] = \text{qr}(\mathbf{Y}, 0)$

### Stage B:

(4) Form the  $(k + p) \times n$  matrix  $\mathbf{B} = \mathbf{Q}^* \mathbf{A}$ .

$\mathbf{B} = \mathbf{Q}' * \mathbf{A}$

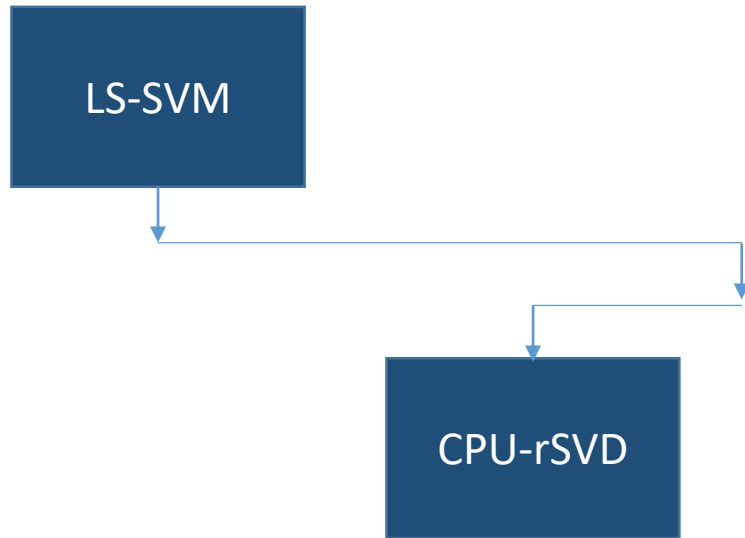
(5) Form the SVD of the small matrix  $\mathbf{B}$ :  $\mathbf{B} = \hat{\mathbf{U}}\mathbf{D}\mathbf{V}^*$ .

$[\mathbf{Uhat}, \mathbf{D}, \mathbf{V}] = \text{svd}(\mathbf{B}, 'econ')$

(6) Form  $\mathbf{U} = \mathbf{Q}\hat{\mathbf{U}}$ .

$\mathbf{U} = \mathbf{Q} * \mathbf{Uhat}$

# Implementation Details

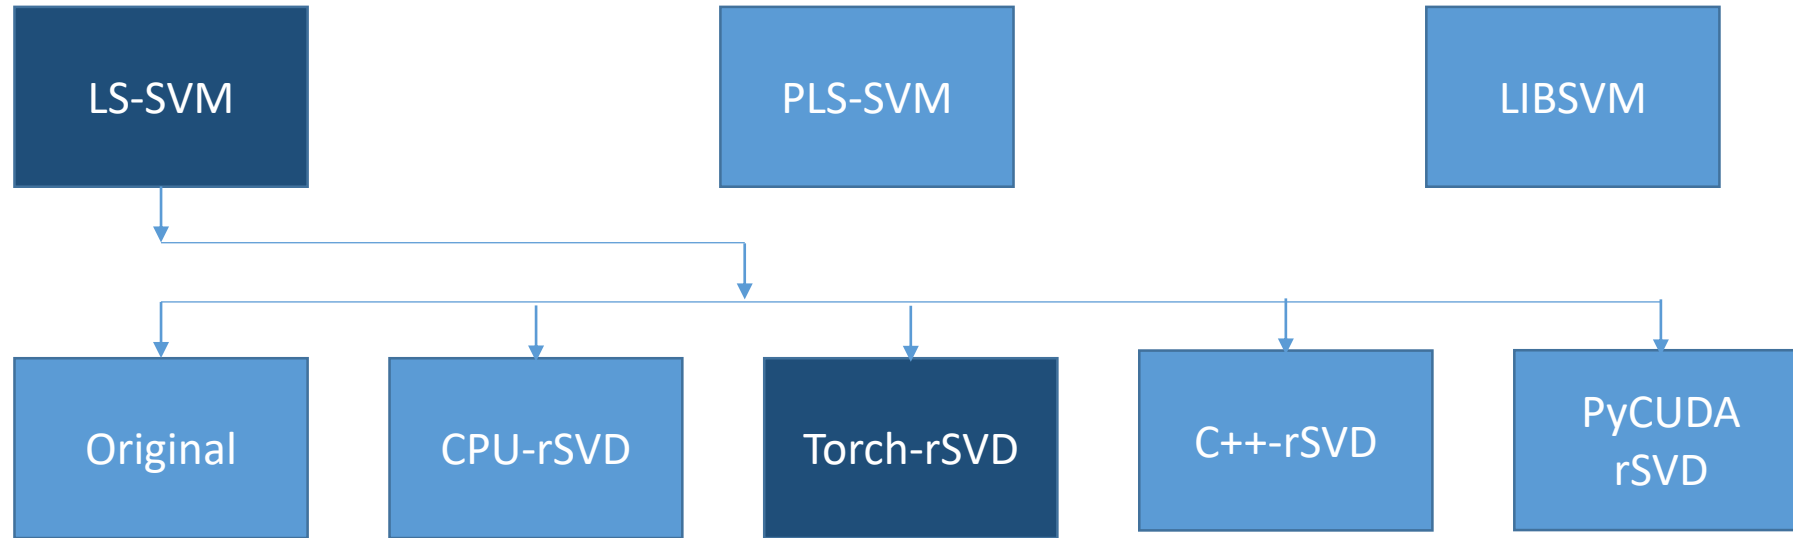


```
if(use_cpu_rsvd):
    print('cpu_rsvd')


    # tic2=time.perf_counter()
    #####
    rank = math.ceil(0.25*Aa.shape[1])
    Omega = np.random.randn(Aa.shape[1], rank)
    Y = Aa @ Omega
    Q, _ = np.linalg.qr(Y)
    Bb = Q.T @ Aa
    u_tilde, s, v = np.linalg.svd(Bb, full_matrices = 0)
    u = Q @ u_tilde
    solution = np.dot(u.T, B)
    solution=np.divide(solution,s)
    solution=np.dot(v.T,solution)
    #####
```

# Implementation Details

---



# Implementation Details



```
if(use_gpu_rsvd):
    print('i D0ing gpu')
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

    at = torch.from_numpy(A)

    at=at.float()
    bb=torch.from_numpy(B)

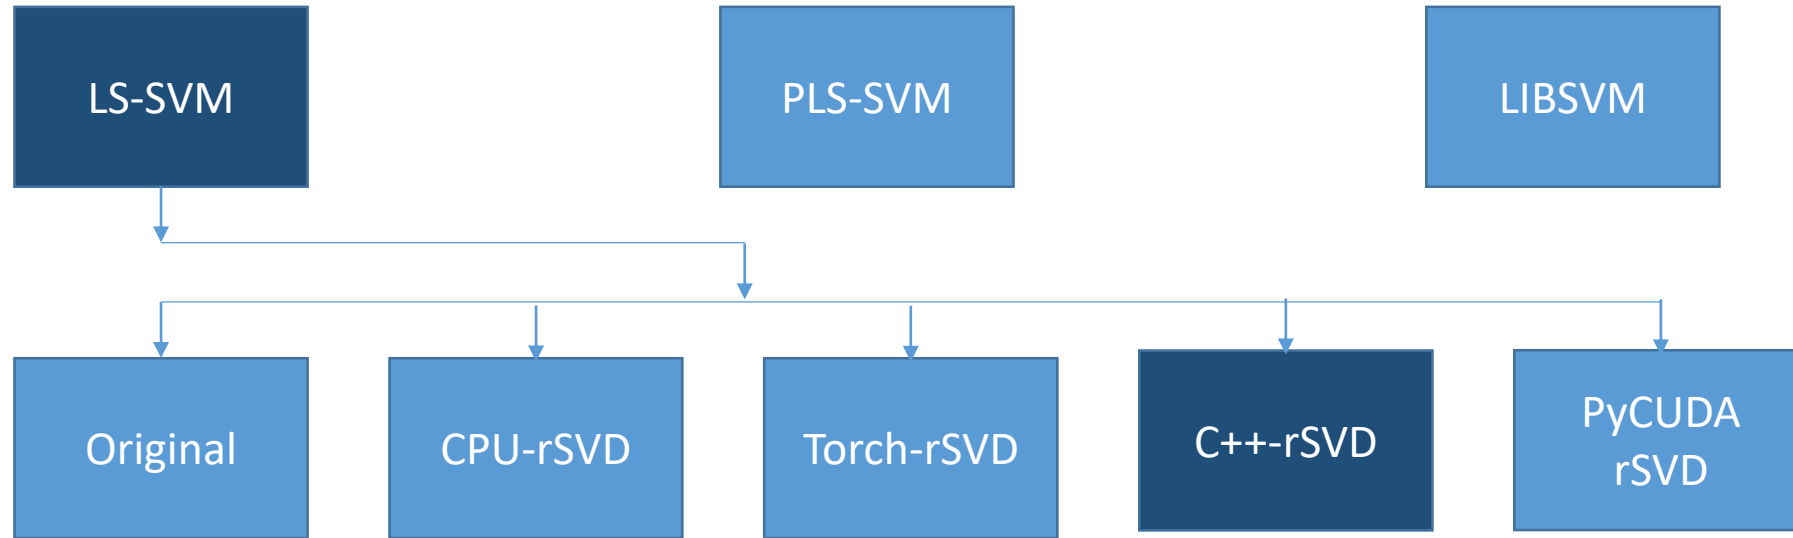
    bb=bb.float()
    tic1=time.perf_counter()
    at=at.cuda()
    bb=bb.cuda()
    with torch.no_grad():

        rank1 = math.ceil(0.75*Aa.shape[1])
        Omegal = torch.rand(at.shape[1],rank1,device=device)
        Yl=torch.matmul(at,Omegal)
        Q1, _ = torch.linalg.qr(at)
        Bb1=torch.matmul(torch.transpose(Q1,0,1),at)
        U_t,S,Vt=torch.linalg.svd(Bb1,full_matrices=False)
        U=torch.matmul(Q1,U_t)
        V=Vt.mH
        solution = torch.mv(U.mH,bb)
        solution = torch.div(solution,S)
        solution = torch.mv(V,solution)

    solution=solution.detach().cpu().numpy()
```

# Implementation Details

---



# Implementation Details

---

LS-SVM



C++-rSVD

```
c302-005.ls6(128)$ gcc -o c_implementation my_mkl.cpp -mkl
```

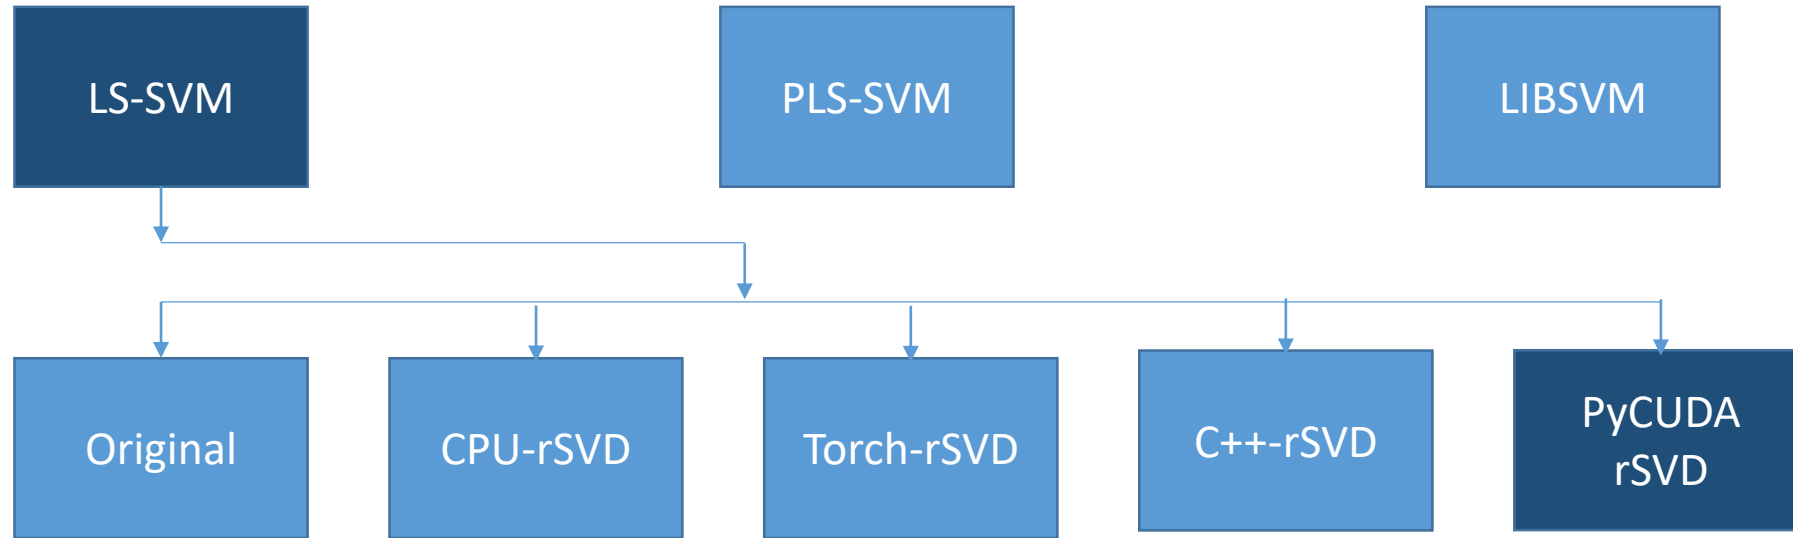
```
cblas_dgemm(CblasRowMajor,CblasNoTrans,CblasNoTrans,m,n,k,1,A,k,0mega,n,0,C,n);
```

```
info = LAPACKE_dgeqrf(LAPACK_ROW_MAJOR,m,n,C,n,tau);
```

```
info = LAPACKE_dorgqr(LAPACK_ROW_MAJOR,m,n,n,C,n,tau);
```

# Implementation Details

---

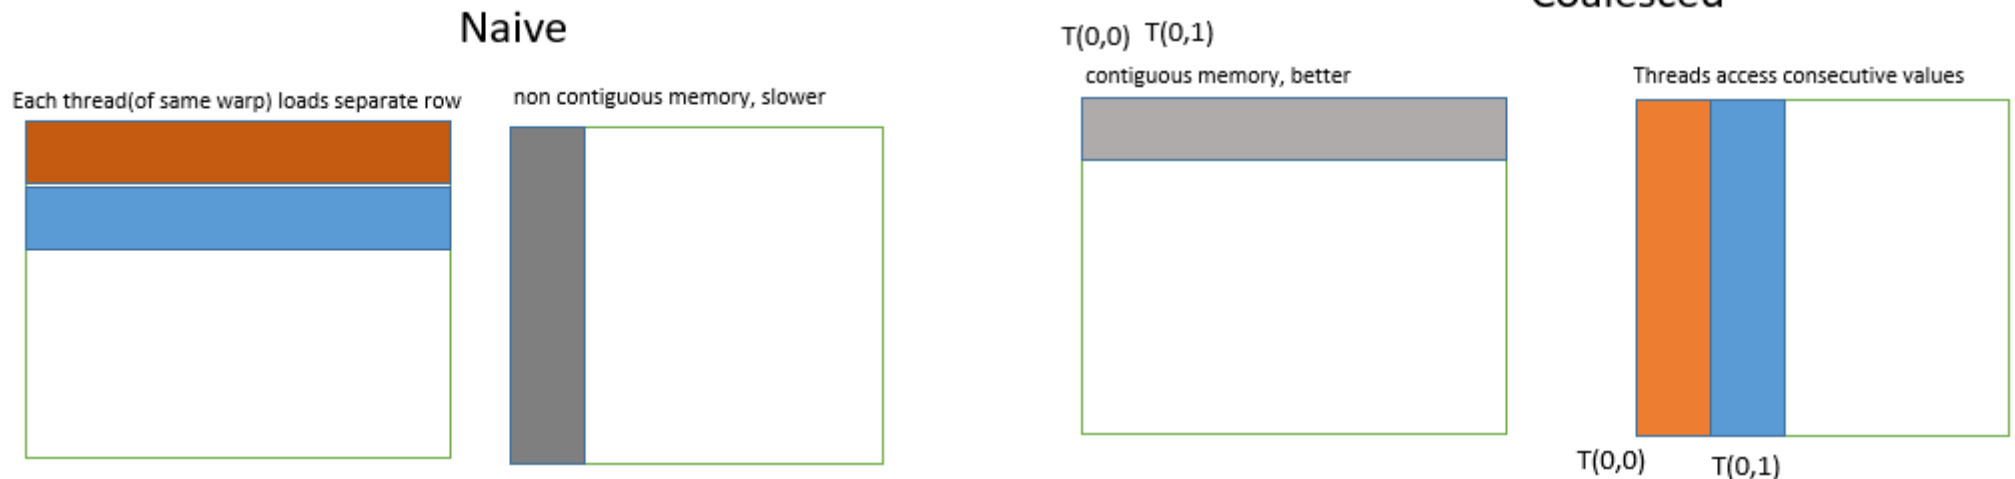



# Implementation Details

## Matrix Multiplication on the GPU:

- Started with the Naïve Kernel. First improvement was by Memory Coalescing the access to data. Threads are launched as a set of warps, with  $T(0,0) \rightarrow T(0,31)$ . But the Naïve implementation makes the threads of the same warp access different memory locations, reducing efficiency. Code changed to reverse access pattern.
- Shared memory used with Tiling.
- Loop unrolling for the lowest level loop.
- Did some benchmark measurements and approximate FLOP counts for this kernel.

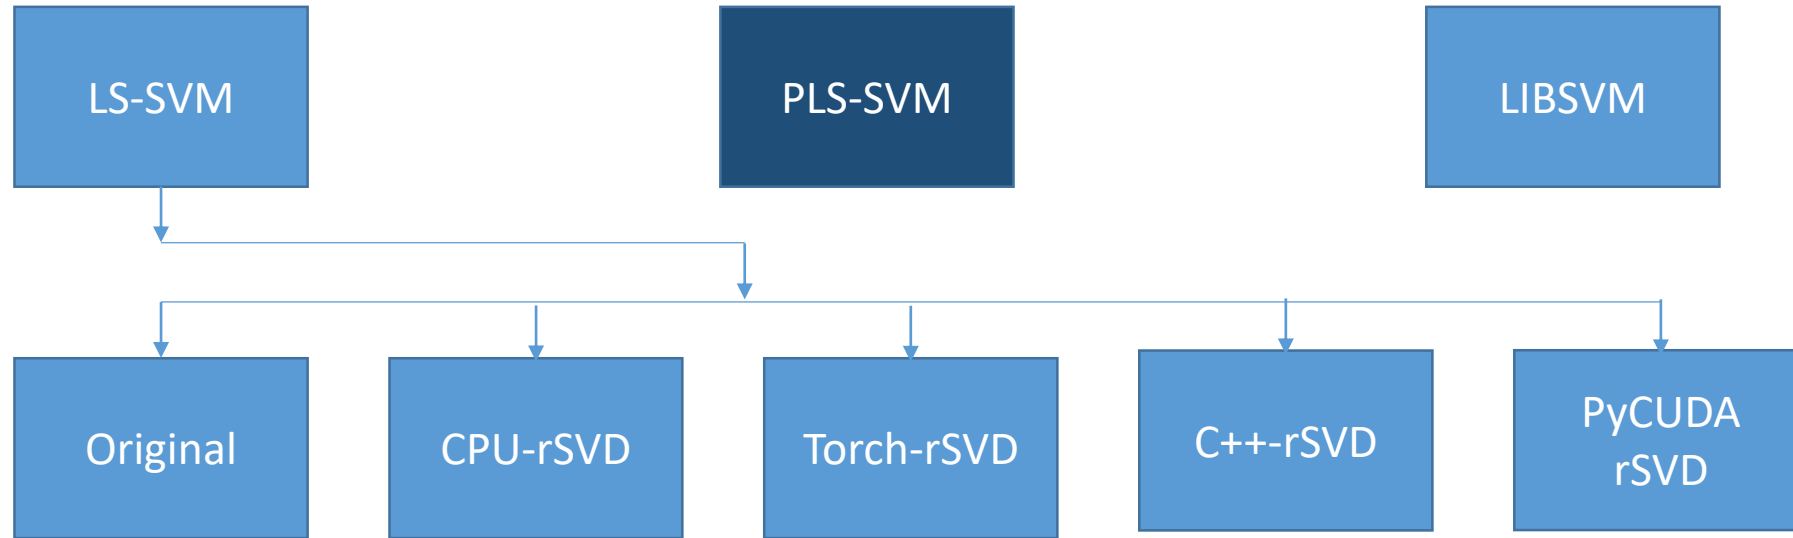
Language: CUDA/C++, also implemented in pycuda.





# Implementation Details

---



# Implementation Details

---

## PLS-SVM

Implementations:

OpenMP  
CUDA  
HIP  
OpenCL  
SYCL

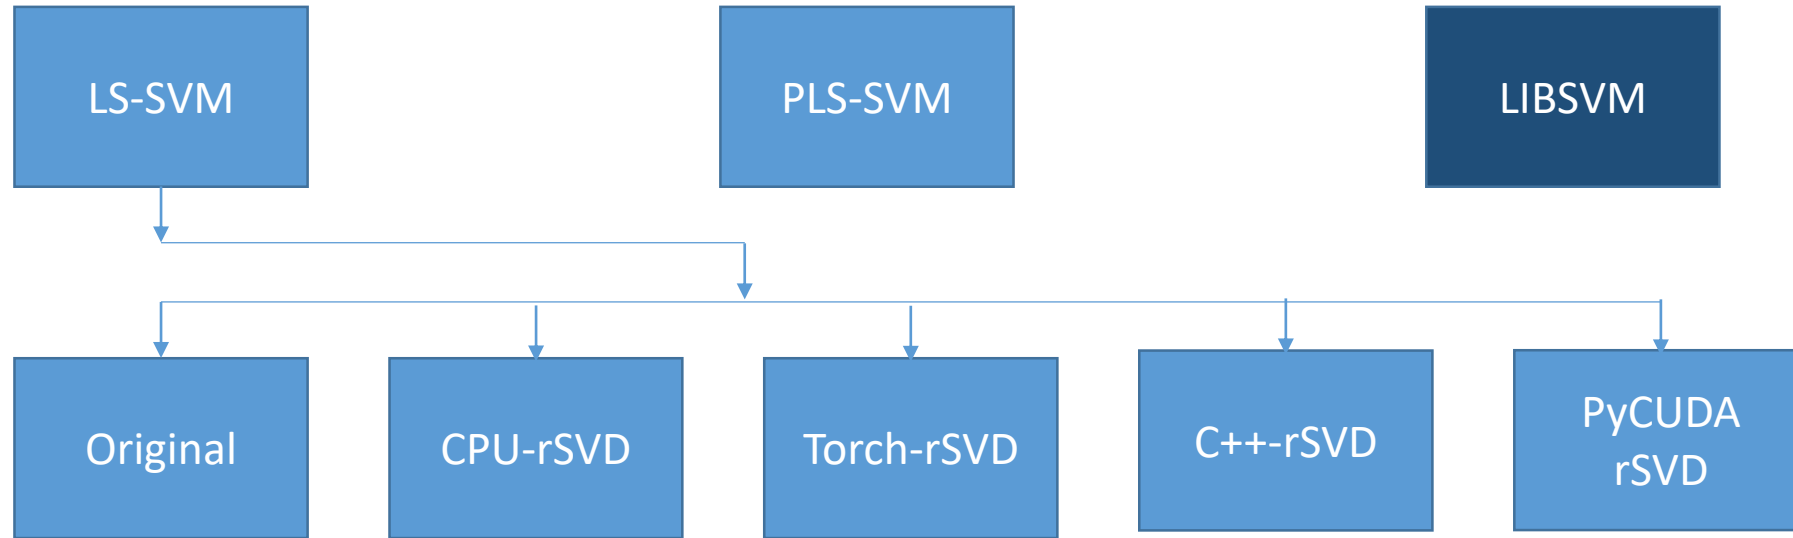
Parallel RSVD:

- Used a simple python implementation of LSSVM\*
- Started with NumPy implementation for the algorithm, used `np.linalg.svd()` and then multiplied them correctly to get the pseudo inverse( $\mathbf{A}^+$ ), and finally did  $\text{solution} = (\mathbf{A}^+ * \mathbf{b})$
- First changed this to the RSVD algorithm and then changed the multiplication order multiply  $\mathbf{b}$  with the `svd(A)` in the correct order to minimize matrix-matrix multiplies.
- To get an idea of the parallelism we could expect, used torch to transfer the data to the GPU, and implemented the RSVD algorithm on the GPU.
- Implemented a C++ BLAS/LAPACKe RSVD algorithm, using blas libraries for QR, svd, etc.
- Implemented a PyCUDA version of RSVD,(using scikit-cuda for QR)where we also plug in our Matrix-Matrix kernel, and measure different benchmarks.

<https://github.com/SC-SGS/PLSSVM>

# Implementation Details

---



# LIBSVM: A Library for Support Vector Machines

Chih-Chung Chang and Chih-Jen Lin

Department of Computer Science

National Taiwan University, Taipei, Taiwan


Email: [cjlin@csie.ntu.edu.tw](mailto:cjlin@csie.ntu.edu.tw)

Initial version: 2001    Last updated: August 23, 2022

`svm-train`, `svm-predict`, `svm-scale`

# Implementation Details

---




# OUTLINE


---

1. Problem Statement
2. Sequential – Complexity & Algorithm
3. Parallel – Complexity & Algorithm
4. Implementation Details
- 5. Results**

# Results



# Results





# OUTLINE

---

1. Problem Statement
2. Sequential – Complexity & Algorithm
3. Parallel – Complexity & Algorithm
4. Implementation Details
5. Results

**END**